




Micro-only latency includes finite HBM, private SRAM ports, compute and cold-start/phase service. Each individual core receives full declared private quotas; shapes have different multiplier counts, so this is not an iso-resource layer comparison or a compute-only timing experiment. Shaded intervals bracket changes in the fastest sampled core; no exact crossover is inferred between discrete Me samples.