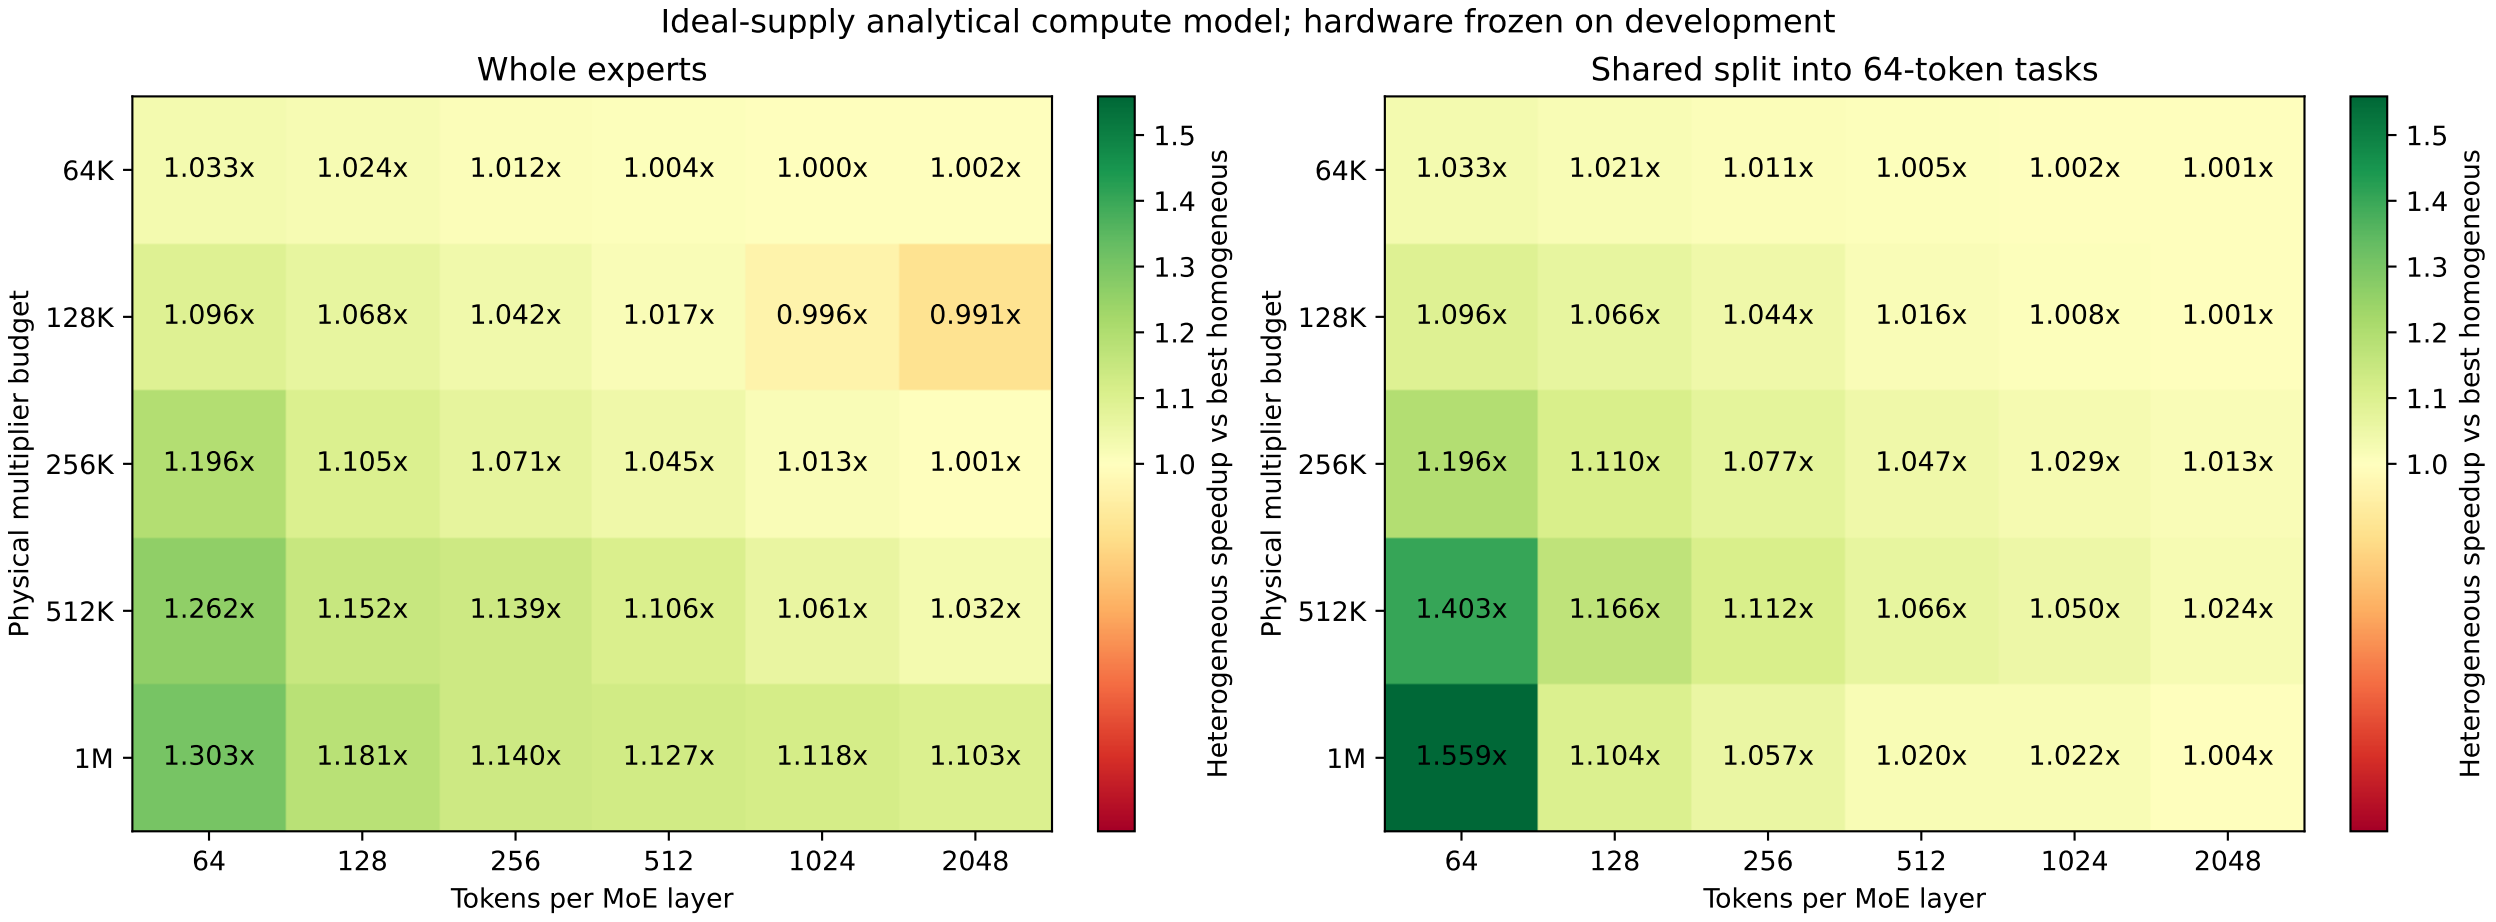
Ideal-supply analytical compute model; hardware frozen on development
Whole experts
Shared split into 64-token tasks
Physical multiplier budget
64K
128K
256K
512K
1M
1.033x 1.024x 1.012x 1.004x 1.000x 1.002x
1.096x 1.068x 1.042x 1.017x 0.996x 0.991x
1.196x 1.105x 1.071x 1.045x 1.013x 1.001x
1.262x 1.152x 1.139x 1.106x 1.061x 1.032x
1.303x 1.181x 1.140x 1.127x 1.118x 1.103x
1.033x 1.021x 1.011x 1.005x 1.002x 1.001x
1.096x 1.066x 1.044x 1.016x 1.008x 1.001x
1.196x 1.110x 1.077x 1.047x 1.029x 1.013x
1.403x 1.166x 1.112x 1.066x 1.050x 1.024x
1.559x 1.104x 1.057x 1.020x 1.022x 1.004x
64 128 256 512 1024 2048
Tokens per MoE layer
Heterogeneous speedup vs best homogeneous
1.5
1.4
1.3
1.2
1.1
1.0